

Regression and Classification Basics

José R. Dorronsoro

Dpto. de Ingeniería Informática, Escuela Politécnica Superior

Universidad Autónoma de Madrid

28049 Madrid, Spain

Machine Learning Modeling Basics

- 1 Machine Learning Modeling Basics
- 2 Basic Linear Regression
- 3 Bias, Variance and Cross Validation
- 4 Basic Classification
- 5 Logistic Regression
- 6 Practical Classification
- 7 Nearest Neighbor Classification

What Is Machine Learning (ML)?

- ▶ Lofty definition: make machines learn!!!
 - ▶ Have to make “machines” and “learn” more precise
- ▶ The machines of ML: mathematical input–output processes that lend themselves to some form of (numerical) parameterization
- ▶ The learning process: adjust the machine’s parameters until a goal is reached
- ▶ New thing: “goal”?
 - ▶ At first sight, get something done
 - ▶ Ultimately, to minimize some error measure
- ▶ Summing things up: a ML process tries to find a concrete mathematical/algorithmic **input–output parameterized transformation** that **minimizes an error measure** by iteratively **adjusting the transformation’s parameters**

Where Lies ML?

- ▶ In the middle of a possibly long process chain
- ▶ Before ML starts we must
 - ▶ Go from **raw to organized** data: accesing, gathering, cleaning, formatting, ...
 - ▶ Go from **organized to** (potentially) **informative** data: extracting basic and derived features
- ▶ After ML finishes and we have a model, we must perform
 - ▶ Outcome **evaluation**: how good/actionable the model is
 - ▶ Outcome **exploitation**: collect, organize, act
 - ▶ **Individual model maintenance**: monitor performance, tune hyper-parameters
 - ▶ **Modeling life cycle maintenance**: discard old models, introduce new ones and **communicate** our work/results
- ▶ ML is in the middle of the global process chain

Supervised/Unsupervised Models

- ▶ First modeling step: get a **sample** $S = \{z^1, \dots, z^N\}$ with $z^p \in \mathbf{R}^D$ the **patterns** and N the sample size
- ▶ Model types: **supervised, unsupervised**
- ▶ Supervised models:
 - ▶ Patterns $z^p = (x^p, y^p)$, with the y^p being the **targets** that the model tries to predict
 - ▶ These known targets guide, or **supervise**, model building
 - ▶ Main emphasis here
- ▶ Unsupervised models:
 - ▶ Patterns $z^p = x^p$ without any known targets
 - ▶ But nevertheless the model is supposed to learn relations or find structure in the data
 - ▶ Sometimes as a first step towards a supervised model

Regression and Classification

- ▶ Problems (usually) to be solved by models: regression, classification
- ▶ Patterns come in pairs (x, y)
 - ▶ x : inputs, predictors, features, independent variables
 - ▶ y : target, response, dependent variable; numerical in regression, class labels in classification
- ▶ **Regression:** the desired output y is regressed into the inputs x to derive a model $\hat{y} = f(x)$
 - ▶ We want $y \simeq \hat{y}$ so having $y - \hat{y}$ “small” is the natural goal
- ▶ **Classification:** inputs are derived from several classes C_1, \dots, C_K , to which labels ℓ_k are assigned
 - ▶ The model now assigns a label $\ell(x)$ to an input x
 - ▶ If x is derived from C_k we want to have $\ell(x) = \ell_k$
 - ▶ Here having $\ell(x) - \ell_k$ “small” may not make sense

How to Build Regression Models

- ▶ In general we have a sample $S = \{x^p, y^p\}$, $1 \leq p \leq N$, with $x^p \in \mathbf{R}^d$ the **features** and y^p the **targets**
- ▶ We want to build a model $\hat{y} = f(x)$ so that $\hat{y}^p = f(x^p) \simeq y^p$; i.e., we want to **regress** y to the x
- ▶ The concrete f is chosen within a certain family \mathcal{F}
 - ▶ Examples here: linear regression, multilayer perceptrons (MLPs), SVMs
 - ▶ And also: Random Forests (RF), Gradient Boosting (GB), Nearest Neighbor (NN)
- ▶ Natural option to ensure $f(x^p) \simeq y^p$: choose f to minimize the sample mean square error (MSE)

$$\hat{e}(f) = \hat{e}_S(f) = \frac{1}{2N} \sum_{p=1}^N (y^p - f(x^p))^2$$

- ▶ Thus, the model we select is $\hat{f} = \hat{f}_S = \arg \min_{f \in \mathcal{F}} \hat{e}_S(f)$

Model Parameterization

- ▶ Usually individual models are selected through (ideally optimal) **parameter sets**
 - ▶ The parameters (weights) $W \in R^M$ select a concrete f in \mathcal{F}
- ▶ **Parametric** models have a fixed functional form $f(x) = f(x; W)$
- ▶ Simplest example: linear regression, where $M = d + 1$ and $W = (w_0, w)$

$$f(x; w_0, w) = w_0 + \sum_{j=1}^d w_j x_j = w_0 + w \cdot x$$

- ▶ **Semi-parametric** models also use weights but without a predefined functional form; MLPs but also RF or GBR
- ▶ **Non parametric** models do not use weights nor follow any broad functional form; Nearest Neighbor models

Model Estimation as Error Minimization

- For a parametric or semiparametric $f(x; W)$ we can write $\widehat{e}_S(f) = \widehat{e}_S(W)$

- The problem to solve becomes

$$\widehat{W}^* = \widehat{W}_S^* = \arg \min_W \widehat{e}_S(f(\cdot; W)), \text{ i.e., } \widehat{e}_S(\widehat{W}^*) \leq \widehat{e}_S(W) \forall W$$

- In linear regression

$$\widehat{e}(w_0, w) = \frac{1}{2N} \sum_{p=1}^N (y^p - w_0 - w \cdot x^p)^2$$

which ends up in a simple **quadratic form**

- The regression problem reduces to **minimize** $\widehat{e}_S(W)$
 - Something in principle well understood in mathematical optimization

Regression Basics

- 1 Machine Learning Modeling Basics
- 2 Basic Linear Regression**
- 3 Bias, Variance and Cross Validation
- 4 Basic Classification
- 5 Logistic Regression
- 6 Practical Classification
- 7 Nearest Neighbor Classification

Linear Models

- ▶ Assuming $x \in \mathbf{R}^d$, the basic linear model is

$$f(x) = w_0 + \sum_1^d w_i x_i = w_0 + w \cdot x$$

- ▶ w_0 complicates notation and, to drop it, we center x and y so that $E[x_i] = E[y] = 0$; then $w_0 = 0$
- ▶ Then we are left with the simpler homogeneous model $f(x) = w \cdot x$
- ▶ In practice we will always **normalize** x , for instance to have 0 mean and 1 standard deviation (std) on each feature
 - ▶ But not y if we may help it
- ▶ But: how do we find w ?

1-dimensional Linear Regression (LR)

- ▶ Assume that features X and target Y are **centered**, i.e., have 0 means
- ▶ For 1-dimensional patterns x the LR model then becomes

$$f(x) = w x$$

- ▶ And the error is then the function $e(w)$

$$\begin{aligned}\widehat{e}(w) &= \frac{1}{2N} \sum_{p=1}^N (w x^p - y^p)^2 = \frac{1}{2N} \sum_p (w^2 (x^p)^2 - 2x^p y^p w + (y^p)^2) \\ &= w^2 \left(\frac{1}{2N} \sum_p (x^p)^2 \right) - w \left(\frac{1}{N} \sum_p x^p y^p \right) + \frac{1}{2N} \sum_p (y^p)^2\end{aligned}$$

- ▶ The problem has obviously a minimum w^*
- ▶ To find it we just solve $\widehat{e}'(w) = 0$

Solving $\widehat{e}'(w) = 0$

- To compute $\widehat{e}'(w)$ we have

$$\widehat{e}'(w) = w \left(\frac{1}{N} \sum_p (x^p)^2 \right) - \left(\frac{1}{N} \sum_p x^p y^p \right)$$

- The optimal w^* solves $\widehat{e}'(w) = 0$ and is given by

$$w^* = \frac{\frac{1}{N} \sum_p x^p y^p}{\frac{1}{N} \sum_p (x^p)^2} = \frac{\frac{1}{N} X \cdot Y}{\frac{1}{N} X \cdot X} = \frac{\frac{1}{N} X \cdot Y}{\text{var}(x)} = \frac{1}{\text{var}(x)} \text{covar}(x, y)$$

where X and Y denote the N dimensional vectors $(x^1, \dots, x^N)^t$, $(y^1, \dots, y^N)^t$

General Linear Regression

- ▶ Assume again that X and Y are centered
- ▶ The LR model becomes now $f(x) = \sum_1^d w_i x_i = w \cdot x$
- ▶ If Y is the $N \times 1$ **target** vector and we organize the sample S in a $N \times d$ **data matrix** X , the sample mse is given by

$$\begin{aligned}\widehat{e}(w) &= \frac{1}{2N} \sum_p (w \cdot x^p - y^p)^2 = \frac{1}{2N} (Xw - Y)^t (Xw - Y) \\ &= \frac{1}{2N} (w^t X^t X w - 2w^t X^t Y + Y^t Y)\end{aligned}$$

- ▶ Now we have to solve $\nabla \widehat{e}(w) = 0$, i.e., $\frac{\partial \widehat{e}}{\partial w_i}(w) = 0$
- ▶ It is easy to see that

$$\nabla \widehat{e}(w) = \frac{1}{N} X^t X w - \frac{1}{N} X^t Y = \widehat{R}w - \widehat{b}$$

Solving the Linear Equations

- ▶ Thus, the optimal \hat{w}^* must solve the **normal equation**
 $\hat{R} \hat{w} - \hat{b} = 0$, where we recall that

$$\hat{R} = \frac{1}{N} X^t X, \quad \hat{b} = \frac{1}{N} X^t Y$$

- ▶ Over the original, non-centered data matrix we have

$$\hat{R} = \frac{1}{N} (X - \bar{X})(X - \bar{X})^t;$$

i.e., \hat{R} is the **sample covariance matrix**

- ▶ If \hat{R} is **invertible**, we just solve the linear system $\hat{R} \hat{w} - \hat{b} = 0$
- ▶ And obtain the sample-dependent optimal \hat{w}^* as

$$\hat{w}^* = \hat{R}^{-1} \hat{b} = (X^t X)^{-1} X^t Y$$

Alternative: Gradient Descent

- ▶ Computing the covariance matrix has a $O(N \times d^2)$ cost and invert it has a $O(d^3)$ cost
 - ▶ For big data problems it may not possible to solve analytically the normal equation $\nabla \hat{e}(w) = 0$
- ▶ The simplest numerical alternative is **gradient descent**:
 - ▶ Starting from some random w^0 we iteratively compute

$$w^{k+1} = w^k - \rho_k \nabla \hat{e}(w^k) = w^k - \frac{\rho}{n_B} \left(\hat{X}_B^t \hat{X}_B w^k - \hat{X}_B^t Y \right)$$

over a **mini-batch** B with n_B samples

- ▶ Component wise: $w_i^{k+1} = w_i^k - \rho_k \frac{\partial \hat{e}}{\partial w_i}(w^k)$
 - ▶ ρ_k is the **learning rate**
- ▶ If $w^k \rightarrow w^*$, then $\nabla \hat{e}(w^*) = 0$
 - ▶ Since our problems have obviously minima, this is enough

Measuring Model Fit

- ▶ First option: **Root Square Error** $RSE = \sqrt{\frac{1}{N} \sum (y^p - \hat{y}^p)^2}$
- ▶ OK, but how good is this? We must always have a **base model** to benchmark our results
- ▶ Simplest “model”: a constant w_0 , which yields the mean $\bar{y} = \frac{1}{N} \sum_1^N y^p$, with square error

$$\frac{1}{N} \sum (y^p - \bar{y})^2 = \text{Var}(y)$$

- ▶ We can compare our model against this base by computing

$$\frac{\sum (y^p - \hat{y}^p)^2}{\sum (y^p - \bar{y})^2} = \frac{RSE^2}{\text{Var}(y)}$$

- ▶ The widely used R^2 coefficient is simply $R^2 = 1 - \frac{RSE^2}{\text{Var}(y)}$

Regularization

- ▶ Our regression solution $\hat{w}^* = (X^t X)^{-1} X^t Y$ won't work if $X^t X$ is not invertible
 - ▶ For instance, when some features are correlated
- ▶ We could fix this working instead with $X^t X + \alpha I$ for some $\alpha > 0$
- ▶ To make this practical, one can show that $\hat{w}^* = (X^t X + \alpha I)^{-1} X^t Y$ minimizes

$$e_R(w) = \frac{1}{2N} \sum_p (y^p - w \cdot x_p^p)^2 + \frac{\alpha}{2} \|w\|^2,$$

- ▶ This is the **Ridge Regression** problem
 - ▶ Our first example of **regularization**, a key technique in Machine Learning
 - ▶ **All ML models must be regularized in some way**
- ▶ Important issue: how to find the right choice for α ?

Takeaways on Linear Regression

1. We introduced **supervised** models
2. We have reviewed the essentials of the **linear regression model** (always the first thing to try)
3. We have considered model estimation as a problem on **error minimization**
4. We have seen how to build linear models **analytically and numerically**
5. We have defined how to **measure model fit**
6. We have introduced **regularization**

Bias, Variance and Cross Validation

- 1 Machine Learning Modeling Basics
- 2 Basic Linear Regression
- 3 Bias, Variance and Cross Validation**
- 4 Basic Classification
- 5 Logistic Regression
- 6 Practical Classification
- 7 Nearest Neighbor Classification

Sample Dependence I

- ▶ **Key assumption:** x and y are related as $y = \phi(x) + n$ where
 - ▶ $\phi(x)$ is the **true** underlying function
 - ▶ n is **additive noise** with 0 mean and finite variance σ_N^2
- ▶ We can't do much about the noise n , so we concentrate in getting a model \hat{f} such that $\hat{f} \simeq \phi(x)$
- ▶ That is, we want our models to have a small **bias**, i.e., get as close as possible to the “true” model ϕ
 - ▶ Intuitively this can be achieved using highly flexible models with many parameters
- ▶ But **everything is sample dependent**: changing S gives a different model
 - ▶ To emphasize this we write $\hat{f}_S(x)$

Sample Dependence II


- ▶ But we also want some model stability with respect to the different samples
- ▶ In other words, for different samples S, S' we want our models to verify

$$\hat{f}_S(x) \simeq \hat{f}_{S'}(x)$$

- ▶ That is, we want our models to have small **variance** with respect to sample changes
 - ▶ Intuitively this can be achieved using simple models with few parameters
- ▶ But obviously the goals of small bias and small variance are contradictory to a large extent

The Bias–Variance Tradeoff

- There is thus a **tradeoff** between bias (low for complex models) and variance (low for simple models)



Taken from *Hastie et al.*, p. 38

Evaluating Expected Performance

- ▶ Samples will change anyway: the sample with which we train our model will be different from the one on which we exploit it
 - ▶ And before we start applying a model, we should have a reasonably accurate idea of its performance in practice
- ▶ To get it over a single S , we apply **Cross Validation** (CV), where we
 - ▶ Randomly split the sample S in M subsets S_1, \dots, S_M
 - ▶ Work with M **folds**: pairs (S_m, S_m^c) , with

$$S_m^c = S - S_m = \cup_{i \neq m} S_i$$

- ▶ Build M different models **using the S_m^c as training subsets**
 - ▶ Compute their errors e_m on the folds' **validation subsets S_m**
 - ▶ Use these **errors' average** as a first estimate of the true model performance
- ▶ CV can and **must be used** in any model building procedure
- ▶ We will also use CV to find an **optimal value** for the hyper-parameter α in Ridge Regression

Grid Hyper-parameter Selection

- ▶ Consider for Ridge regression a hyper-parameter range $[0, A]$
 - ▶ $\alpha = 0$: no penalty and, thus, small bias but possibly high variance
 - ▶ $\alpha = A$: large penalty and, thus, small variance but high bias
- ▶ Select an $L + 1$ point **grid** $[\alpha_0, \dots, \alpha_L]$
 - ▶ For instance a uniform one $\alpha_\ell = \ell \frac{A}{L}, \ell = 0, 1, \dots, L$
- ▶ Build M **folds**: pairs (S_m, S_m^c) and for each α_ℓ
 - ▶ Train M Ridge models on the S_m^c using the hyper-parameter α_ℓ
 - ▶ Average their M validation errors e_m on the S_m to get the CV error $e(\alpha_\ell)$ for α_ℓ
- ▶ Finally choose the (hopefully) optimal hyper-parameter α^* as

$$\alpha^* = \arg \min_{0 \leq \ell \leq L} e(\alpha_\ell)$$

- ▶ α^* gives the model with the **best expected generalization among all possible α choices**

Takeaways on Bias, Variance and CV

1. We have stressed that **any model estimation is sample-dependent** and that this has to be controlled
2. We have introduced the **bias** and **variance** as the two key components of any model error
3. We have discussed **bias-variance trade-off**
4. We have introduced **Cross Validation** here as a tool to estimate a **model's generalization performance**
5. We have also introduced **Cross Validation** as a tool to estimate a **model's hyper-parameters**

Basic Classification

- 1 Machine Learning Modeling Basics
- 2 Basic Linear Regression
- 3 Bias, Variance and Cross Validation
- 4 Basic Classification**
- 5 Logistic Regression
- 6 Practical Classification
- 7 Nearest Neighbor Classification

Regression vs Classification

- ▶ Recall that in regression we have numerical continuous targets y and want our predictions \hat{y} to be as close to y as possible
- ▶ But in classification we have a finite number of labelled targets for which “selection by closeness” doesn’t make sense
- ▶ Natural alternative: select the **most probable label given the pattern** we have just received
 - ▶ The concrete labels used for targets do not matter anymore
 - ▶ Model learning should thus be “target” agnostic
 - ▶ And good probability estimates should be quite useful
- ▶ Let’s analyze this in an example

A First Problem: Pima Indian Diabetes

- ▶ We want to diagnose whether a person may have diabetes from some clinical measures
- ▶ Features x : clinical measures
 - ▶ numPregnant
 - ▶ bloodPress
 - ▶ massIndex
 - ▶ age ...
- ▶ Target y : 0 (no diabetes), 1 (diabetes)
- ▶ Clear goal but perhaps too radical
- ▶ Better: try to estimate **the probability $P(1|x)$ of having diabetes depending on the features x we measure**

Classification Setup

- ▶ We have random patterns ω from M classes, C_1, \dots, C_M
- ▶ Over each pattern we “measure” d features $x = x(\omega) \in \mathbb{R}^d$
 - ▶ x inherits the randomness in ω and becomes a **random variable**
- ▶ A ω has a **prior probability** π_m of belonging to C_m
- ▶ Inside each class C_m there is a **conditional class density** $f(x|m)$ that “controls” the appearance of a given x
- ▶ The π_m and $f(x|m)$ determine the **posterior probability** $P(m|x)$ that x comes from class C_m
- ▶ **Intuition:** we should assign x to the class with the **largest** $P(m|x)$, that is, work with the classifier

$$\delta(x) = \arg \max_m P(m|x)$$

The Obviously Optimal Classifier

- ▶ It can be shown that $P(m|x) = \frac{\pi_m f(x|m)}{f(x)}$
- ▶ Thus, we should decide according to a **classifier** function δ_B

$$\begin{aligned}\delta_B(x) &= \arg \max_m P(m|x) = \arg \max_m \frac{\pi_m f(x|m)}{f(x)} \\ &= \arg \max_m \pi_m f(x|m)\end{aligned}$$

- ▶ With some extra work we can show that this **Bayes Classifier** δ_B defines an optimal solution (in some precise sense) of the classification problem
- ▶ But ... this doesn't look too practical for we do not know either π_m or (much harder) $f(x|m)$
- ▶ We will focus first on estimating **directly** the label generating distribution $P(m|x)$

Logistic Regression (LR)

- 1 Machine Learning Modeling Basics
- 2 Basic Linear Regression
- 3 Bias, Variance and Cross Validation
- 4 Basic Classification
- 5 Logistic Regression**
- 6 Practical Classification
- 7 Nearest Neighbor Classification

Logistic Regression (LR)

- ▶ We assume

$$P(1|x) = P(1|x; w_0, w) = \frac{1}{1 + e^{-(w_0 + w \cdot x)}}$$

- ▶ Then $0 \leq P(1|x) \leq 1$ for any x
- ▶ We then have

$$P(0|x) = 1 - P(1|x) = \frac{e^{-(w_0 + w \cdot x)}}{1 + e^{-(w_0 + w \cdot x)}} = \frac{1}{1 + e^{w_0 + w \cdot x}}$$

- ▶ Notice that if $w_0 + w \cdot x = 0$, $P(1|x) = P(0|x) = 0.5$
- ▶ The ratio $\frac{P(1|x)}{P(0|x)} = e^{w_0 + w \cdot x}$ is called the **odds** of x and its log the **log odds** or **logit**
- ▶ Thus, the basic assumption in LR is that the **logit is a linear function** $w_0 + w \cdot x$ of x
- ▶ We have the model $f(x; w)$; we need a **loss function** $L(w_0, w)$ to minimize for which we use the sample's **likelihood**

Estimating w_0^*, w^*

- ▶ Assume a single sample x, y and two possible model coefficients w_0, w and w'_0, w'
- ▶ Denoting by $p = P(y|x; w_0, w)$ and $p' = P(y|x; w'_0, w')$, it is clear that we should prefer w_0, w if $p > p'$ and w'_0, w' if not
 - ▶ In other words, we prefer the coefficients that give a **higher posterior probability** to the sample (x, y)
- ▶ For an independent sample $S = \{(x^p, y^p)\}$, its joint probability under a posterior model $p = P(y|x; w_0, w)$ is

$$P(Y|X; w_0, w) = \prod_{p=1}^N P(y^p|x^p; w_0, w)$$

- ▶ And, again, given two possible model coefficients w_0, w and w'_0, w' , we should prefer w_0, w iff

$$P(Y|X; w_0, w) > P(Y|X; w'_0, w')$$

Sample's Likelihood

- Therefore, we can estimate the optimal w_0^*, w^* as

$$w_0^*, w^* = \arg \max_{w_0, w} P(Y|X; w_0, w)$$

- By the independence assumption we have

$$\begin{aligned} P(Y|X; w_0, w) &= \left\{ \prod_{y^p=1} P(1|x^p) \right\} \left\{ \prod_{y^p=0} P(0|x^p) \right\} \\ &= \prod_{p=1}^N P(1|x^p)^{y^p} P(0|x^p)^{1-y^p} \end{aligned}$$

with the last equality follows from

- If $y^p = 1$, $P(1|x^p) = P(1|x^p)^{y^p} = P(1|x^p)^{y^p} P(0|x^p)^{1-y^p}$, and
- If $y^p = 0$, $P(0|x^p) = P(0|x^p)^{1-y^p} = P(1|x^p)^{y^p} P(0|x^p)^{1-y^p}$

Max Log-Likelihood Estimation

- The log-likelihood of w_0, w given S is then

$$\begin{aligned}\ell(w_0, w; S) &= \log P(Y|X; w_0, w) \\ &= \sum_{p=1}^N \{y^p \log P(1|x^p) + (1 - y^p) \log P(0|x^p)\} \\ &= \sum_p y^p \log \frac{P(1|x^p)}{P(0|x^p)} + \sum_p \log P(0|x^p) \\ &= \sum_p y^p (w_0 + w \cdot x^p) - \sum_p \log(1 + e^{w_0 + w \cdot x^p})\end{aligned}$$

- We can thus estimate the optimal \hat{w}_0^*, \hat{w}^* as

$$\hat{w}_0^*, \hat{w}^* = \arg \min_{w_0, w} -\ell(w_0, w; S)$$

- Extra bonus: $-\ell$ is a convex differentiable function of (w_0, w) and, thus, it is enough to solve $\nabla \ell(w_0, w) = 0$

Multinomial Logistic Regression

- ▶ It is applied to multiclass problems
- ▶ Assuming K classes, the posterior probabilities are assumed to be

$$P(c|x) = \frac{e^{w_0^c + w^c \cdot x}}{1 + \sum_1^{K-1} e^{w_0^k + w^k \cdot x}}, \quad 1 \leq c \leq K-1$$
$$P(K|x) = \frac{1}{1 + \sum_1^{K-1} e^{w_0^k + w^k \cdot x}}$$

- ▶ Given a sample S , we can define a variant $\ell(W_0, W)$ of the log likelihood approach followed in logistic regression
 - ▶ Here $W_0 \in \mathbf{R}^K$ and W is a $K \times d$ matrix, with d the x dimension
- ▶ Again, we can estimate the optimal \hat{W}_0^*, \hat{W}^* as

$$\hat{W}_0^*, \hat{W}^* = \arg \min_{W_0, W} -\ell(W_0, W; S)$$

- ▶ Again $-\ell$ is a convex differentiable function and we just have to solve $\nabla \ell(W_0, W) = 0$

Newton–Raphson Solution

- ▶ However, neither in logistic nor in multinomial regression does $\nabla \ell(\mathcal{W}) = 0$ admit a closed form solution but only an iterative, numerical one
 - ▶ Here we write \mathcal{W} for either (w_0, w) or (W_0, W)
- ▶ We solve it with the **Newton–Raphson** iterative method (equivalent here to Newton’s method for minimization)
- ▶ Starting from a random \mathcal{W}^0 , Newton’s iterations are

$$\mathcal{W}^{k+1} = \mathcal{W}^k + (\mathcal{H}_\ell(\mathcal{W}^k))^{-1} \nabla \ell(\mathcal{W}^k)$$

- ▶ $\mathcal{H}_\ell(\mathcal{W}^k)$ denotes the Hessian of ℓ at \mathcal{W}^k (which may or may not be invertible)
 - ▶ Everything is fine if the \mathcal{W}^k are close enough to the optimum \mathcal{W}^* but far away things may get tricky
- ▶ Just as before, we can add regularization terms $\frac{\alpha}{2} \|w\|^2$ or $\frac{\alpha}{2} \|W\|^2$
- ▶ The iterations in Logistic and Multinomial Regression are again typical of many of the model building methods used in Machine

Learning in ML

- ▶ The general approach to **learning** is usually the following:
 - ▶ A **model** $f(x; W)$ is chosen
 - ▶ Given a sample $S = \{(x^1, y^1), \dots, (x^N, y^N)\}$ and a loss function $\ell(y, \hat{y})$, we define a **sample dependent loss function**

$$L(W) = L(W|S) = \sum \ell(y^p, \hat{y}^p = f(x^p; W))$$

- ▶ $L(W)$ is often minimized from some W^0 by **iterations**

$$W^{k+1} = W^k - \rho_k G(W^k, S)$$

with ρ_k a **learning rate** and G some vectorial function

- ▶ When $G(W) = \nabla L(W)$ we have **gradient descent**
 - ▶ When $G(W) = \mathcal{H}(W)^{-1} \nabla L(W)$ we obtain **Newton's method**
- ▶ In **batch learning** the entire sample S is used at each iteration
- ▶ **On-line** or **minibatch learning**: we use either a single patterns (x^p, y^p) or small subsample
- ▶ Several such procedures will appear here in the coming weeks

Practical Classification

- 1 Machine Learning Modeling Basics
- 2 Basic Linear Regression
- 3 Bias, Variance and Cross Validation
- 4 Basic Classification
- 5 Logistic Regression
- 6 Practical Classification**
- 7 Nearest Neighbor Classification

True/False Positives/Negatives

- ▶ Consider a two class problem with labels $y = 0, 1$
- ▶ We will call patterns with label 1 **positive** and those with label 0 **negative**
 - ▶ Usually the positive patterns are the interesting ones: sick people, defaulted loans, ...
- ▶ Let $\hat{y} = \hat{y}(x)$ the label predicted at x ; we say that x is a
 - ▶ **True Positive (TP)** if $y = \hat{y} = 1$
 - ▶ **True Negative (TN)** if $y = \hat{y} = 0$
 - ▶ **False Positive (FP)** if $y = 0$ but $\hat{y} = 1$
 - ▶ **False Negative (FN)** if $y = 1$ but $\hat{y} = 0$
- ▶ The standard way of presenting these data is through the **confusion matrix**

The Confusion Matrix

- Standard layout

	P' (Predicted)	N' (Predicted)
P (Actual)	True Positive	False Negative
N (Actual)	False Positive	True Negative

- Other layouts:

- Positives(with label 1) at bottom, as done in `confusion_matrix` of `sklearn`
- Predicted values in rows, real values in columns

Classifier Metrics

- ▶ The classifier **accuracy** is $acc = \frac{TP+TN}{N}$
- ▶ acc is the first thing to measure but it may not be too significant: if the number N_0 of negatives is $\gg N_1$, the number of positives
 - ▶ The classifier $\delta(x) = 0$ will have a high accuracy $N_0/N \simeq 1$
 - ▶ But it will also be useless!!
- ▶ First variant in two class problems: Precision, Recall
 - ▶ **Recall** or **sensitivity** or **true positive rate** : $TP/(TP + FN)$, i.e., the fraction of positives detected
 - ▶ **Precision** or **positive predictive value**: $TP/(TP + FP)$, i.e., the fraction of true alarms issued
 - ▶ Also **specificity** or **true negative rate**: $TN/(TN + FP)$,
- ▶ Recall measures how many positive cases we recover, i.e., how **effective** is our method
- ▶ Precision measures the effort we need for that, i.e., its **efficiency**
- ▶ Ideal classifier: high recall, high precision (i.e., effective and efficient!!)

How to Handle Posterior Probabilities

- ▶ If possible, we want **posterior probabilities** as model outputs better than labels
- ▶ Most models give them as pairs

$$(\hat{P}(0|x), \hat{P}(1|x)) = (\hat{P}(0|x), 1 - \hat{P}(0|x))$$

- ▶ In principle we would decide 1 if $\hat{P}(1|x) > 0.5$ and viceversa, but this may be too crude
- ▶ It may be advisable to set a **decision threshold** κ and decide 1 if $\hat{P}(1|x) > \kappa$ and 0 if $\hat{P}(1|x) < \kappa$
- ▶ For **imbalanced** problems where $\pi_0 \gg \pi_1$ (usually the interesting ones) we would have $\hat{P}(1|x) \simeq 0$ for most x
 - ▶ In this case we may choose a $\kappa < 0.5$ and **suggest** 1 if $\hat{P}(1|x) > \kappa$

What's New from Regression?

- ▶ Some things change from regression, some don't
- ▶ We should check feature correlations, if only to remove too similar features
- ▶ Important: **positive and negative-class feature histograms**
 - ▶ Scatter plots (x_i, y) are usually less informative
- ▶ The **bias-variance trade-off** is subtler in classification
- ▶ Accuracy, recall, precision are the usual model quality measures
- ▶ We use CV with **stratified folds** to estimate generalization performance
- ▶ We also use CV for hyperparameter estimation, as regularization will also be needed
 - ▶ In LR we should minimize $-\ell(w_0, w; S) + \frac{\alpha}{2} \|w\|^2$

Takeaways on Basic Classification

1. We have introduced the classification problem as one of computing **posterior probabilities**
2. We have defined the **optimal Bayes classifier**
3. We have introduced **Logistic Regression** to estimate posterior probabilities, and the numerical minimization of its (minus) log-likelihood
4. We have introduced **accuracy, recall, precision** as first classification metrics
 - ▶ Accuracy also works for multiclass problems; recall or precision do not extend so well
5. We have reviewed some practical issues in classification

Nearest Neighbor Classification

- 1 Machine Learning Modeling Basics
- 2 Basic Linear Regression
- 3 Bias, Variance and Cross Validation
- 4 Basic Classification
- 5 Logistic Regression
- 6 Practical Classification
- 7 Nearest Neighbor Classification**

Approximating the Bayes Classifier

- ▶ Starting with the approximation $P(m|x) \simeq P(m|B_r(x))$, Bayes formula gives

$$P(m|x) \simeq P(m|B_r(x)) = \frac{P(C_m \cap B_r(x))}{P(B_r(x))} = \frac{\pi_m P(B_r(x)|m)}{P(B_r(x))}$$

- ▶ Fix x and k and assume a N pattern sample with
 - ▶ N_m patterns are in C_m
 - ▶ $B(x, r)$ the smallest ball containing k sample patterns
 - ▶ $B(x, r)$ has k_m patterns from class C_m
- ▶ We then have the following approximations

$$\pi_m \simeq \frac{N_m}{N}, \quad P(B_r(x)) \simeq \frac{k}{N}, \quad P(B_r(x)|m) \simeq \frac{k_m}{N_m}$$

The k -NN Classifier

- We can thus approximate $P(m|x)$ as

$$P(m|x) \simeq \frac{\pi_m P(B_r(x)|m)}{P(B_r(x))} \simeq \frac{k_m}{N_m} \frac{N_m}{N} \frac{1}{\frac{k}{N}} = \frac{k_m}{k}$$

- And the optimal δ_B as

$$\delta_B(x) \simeq \arg \max_m P(m|B_r(x)) = \arg \max_m \frac{k_m}{k}$$

- We have thus arrived to the k -**Nearest Neighbor** classifier

$$\delta_k^{NN}(x) = \arg \max_m \frac{k_m}{k} = \arg \max_m k_m$$

- Thus $\delta_{kNN}(x)$ assigns x to the class that **has more patterns in $N_k(x)$, the subset of the k neighbors that are closest to x**

By the Way: k -NN Regression

- ▶ k -NN Classification assumes that a pattern should belong to the class with the closest patterns
- ▶ k -NN Regression also relies on a similar assumption: **Predictors that are close should give predictions that are also close**
- ▶ In k -NN Regression we fix a number k of neighbors to be considered and for an input x set

$$\hat{y} = Y_{kNN}(x) = \frac{1}{k} \sum_{x^p \in N_k(x)} y^p$$

where $N_k(x)$ denotes again the k sample points closest to x

- ▶ **Weighted variants:** for instance,
$$Y_k^w(x) = \frac{1}{C_k(x)} \sum_{x^p \in N_k(x)} \frac{1}{\|x^p - x\|} y^p$$
 - ▶ $C_k(x) = \sum_{x^p \in N_k(x)} \frac{1}{\|x^p - x\|}$ is a normalizing constant

Some k -NN Issues

- ▶ **Q1: How do we choose k ?** Using CV, of course, and balancing again the bias–variance tradeoff
 - ▶ Small variance with large k : if $k = N$, k -NN regression returns the mean
 - ▶ Small bias with small k : if $k = 1$ a very close point should give a very close prediction
 - ▶ But also large variance: the nearest point to x in another sample may have a quite different target or belong to another class
- ▶ **Q2: Is k -NN meaningful?** Yes but only if predictors that are close give predictions that are also close, and **provided that there are enough of them close by**
- ▶ This will not be easy because of the **curse of dimensionality**:
 - ▶ Even for low dimensions and large samples, **the sample space is essentially empty**
 - ▶ And for most problems, **there never will be close enough points**
 - ▶ Thus, to get k observations we may go too far away from x and the k -NN predictions will not be meaningful